


```

In [1]: import math
import numpy as np
import time
from __future__ import print_function
import sys
import tensorflow as tf
from collections import deque
import shutil
import random
import os
sys.setrecursionlimit(5000)
class dotdict(dict):
    def __getattr__(self, name):
        return self[name]

args = {
    'numIters': 1000,
    'numEps': 100,
    'tempThreshold': 15,
    'updateThreshold': 0.6,
    'maxlenOfQueue': 200000,
    'numMCTSSims': 25,
    'arenaCompare': 40,
    'cpuct': 1,

    'checkpoint': './temp/',
    'load_model': False,
    'load_folder_file': ('/models', 'best.pth.tar'),
    'lr': 0.001,
    'dropout': 0.3,
    'epochs': 10,
    'batch_size': 64,
    'num_channels': 512,
}

class MCTS():
    """
    This class handles the MCTS tree.
    """

    def __init__(self, game, nnet, args):
        self.game = game
        self.nnet = nnet
        self.args = args
        self.Qsa = {}          # stores Q values for s,a (as defined in the pap
        self.Nsa = {}          # stores #times edge s,a was visited
        self.Ns = {}           # stores #times board s was visited
        self.Ps = {}           # stores initial policy (returned by neural net)

        self.Es = {}          # stores game.getGameEnded ended for board s
        self.Vs = {}          # stores game.getValidMoves for board s

    def getActionProb(self, canonicalBoard, temp=1):
        """
        This function performs numMCTSSims simulations of MCTS starting from
        canonicalBoard.

```

```

Returns:
    probs: a policy vector where the probability of the ith action is
           proportional to Nsa[(s,a)]**(1./temp)
    """
    for i in range(self.args['numMCTSSims']):
        self.search(canonicalBoard)

    s = self.game.stringRepresentation(canonicalBoard)
    counts = [self.Nsa[(s,a)] if (s,a) in self.Nsa else 0 for a in range

    if temp==0:
        bestA = np.argmax(counts)
        probs = [0]*len(counts)
        probs[bestA]=1
        return probs

    counts = [x**(1./temp) for x in counts]
    probs = [x/float(sum(counts)) for x in counts]
    return probs

```

```

def search(self, canonicalBoard):
    """

```

This function performs one iteration of MCTS. It is recursively called till a leaf node is found. The action chosen at each node is one that has the maximum upper confidence bound as in the paper.

Once a leaf node is found, the neural network is called to return an initial policy P and a value v for the state. This value is propagated up the search path. In case the leaf node is a terminal state, the outcome is propagated up the search path. The values of Ns, Nsa, Qsa are updated.

NOTE: the return values are the negative of the value of the current state. This is done since v is in [-1,1] and if v is the value of a state for the current player, then its value is -v for the other player.

```

Returns:
    v: the negative of the value of the current canonicalBoard
    """

```

```

s = self.game.stringRepresentation(canonicalBoard)

if s not in self.Es:
    self.Es[s] = self.game.getGameEnded(canonicalBoard, 1)
if self.Es[s]!=0:
    # terminal node
    return -self.Es[s]

if s not in self.Ps:
    # leaf node
    self.Ps[s], v = self.nnet.predict(canonicalBoard)
    valids = self.game.getValidMoves(canonicalBoard, 1)
    self.Ps[s] = self.Ps[s]*valids          # masking invalid moves
    self.Ps[s] /= np.sum(self.Ps[s])      # renormalize

```

```

        self.Vs[s] = valids
        self.Ns[s] = 0
        return -v

valids = self.Vs[s]
cur_best = -float('inf')
best_act = -1

# pick the action with the highest upper confidence bound
for a in range(self.game.getActionSize()):
    if valids[a]:
        if (s,a) in self.Qsa:
            u = self.Qsa[(s,a)] + self.args['cpuct']*self.Ps[s][a]*m
        else:
            u = self.args['cpuct']*self.Ps[s][a]*math.sqrt(self.Ns[s]

        if u > cur_best:
            cur_best = u
            best_act = a

a = best_act
next_s, next_player = self.game.getNextState(canonicalBoard, 1, a)
next_s = self.game.getCanonicalForm(next_s, next_player)

v = self.search(next_s)

if (s,a) in self.Qsa:
    self.Qsa[(s,a)] = (self.Nsa[(s,a)]*self.Qsa[(s,a)] + v)/(self.Ns[
    self.Nsa[(s,a)] += 1

else:
    self.Qsa[(s,a)] = v
    self.Nsa[(s,a)] = 1

self.Ns[s] += 1
return -v

```

```

class Arena():

```

```

    """

```

```

    An Arena class where any 2 agents can be pit against each other.

```

```

    """

```

```

    def __init__(self, player1, player2, game, display=None):

```

```

        """

```

```

        Input:

```

```

            player 1,2: two functions that takes board as input, return action

```

```

            game: Game object

```

```

            display: a function that takes board as input and prints it (e.g
                    display in othello/OthelloGame). Is necessary for verbose
                    mode.

```

```

        see othello/OthelloPlayers.py for an example. See pit.py for pitting
        human players/other baselines with each other.

```

```

        """

```

```

        self.player1 = player1

```

```

        self.player2 = player2

```

```

        self.game = game

```

```

        self.display = display

def playGame(self, verbose=False):
    """
    Executes one episode of a game.

    Returns:
        either
            winner: player who won the game (1 if player1, -1 if player2
        or
            draw result returned from the game that is neither 1, -1, no
    """
    players = [self.player2, None, self.player1]
    curPlayer = 1
    board = self.game.getInitBoard()
    it = 0
    while self.game.getGameEnded(board, curPlayer)==0:
        it+=1
        if verbose:
            assert(self.display)
            print("Turn ", str(it), "Player ", str(curPlayer))
            self.display(board)
        action = players[curPlayer+1](self.game.getCanonicalForm(board,
        curPlayer))
        valids = self.game.getValidMoves(self.game.getCanonicalForm(board,
        curPlayer))
        if valids[action]==0:
            print(action)
            assert valids[action] >0
        board, curPlayer = self.game.getNextState(board, curPlayer, action)
    if verbose:
        assert(self.display)
        print("Game over: Turn ", str(it), "Result ", str(self.game.getGameEnded(board, curPlayer)))
        self.display(board)
    return self.game.getGameEnded(board, curPlayer)

def playGames(self, num, verbose=False):
    """
    Plays num games in which player1 starts num/2 games and player2 starts
    num/2 games.

    Returns:
        oneWon: games won by player1
        twoWon: games won by player2
        draws: games won by nobody
    """
    #
    end = time.time()
    print('in playGames')
    eps = 0
    maxeps = int(num)

    num = int(num/2)
    oneWon = 0
    twoWon = 0
    draws = 0
    for _ in range(num):
        gameResult = self.playGame(verbose=verbose)

```

```

        if gameResult==1:
            oneWon+=1
        elif gameResult==-1:
            twoWon+=1
        else:
            draws+=1
        # bookkeeping + plot progress
        eps += 1
        #eps_time.update(time.time() - end)
#         end = time.time()

self.player1, self.player2 = self.player2, self.player1

for _ in range(num):
    gameResult = self.playGame(verbose=verbose)
    if gameResult==-1:
        oneWon+=1
    elif gameResult==1:
        twoWon+=1
    else:
        draws+=1
    # bookkeeping + plot progress
    eps += 1
#     eps_time.update(time.time() - end)
    end = time.time()

    return oneWon, twoWon, draws
class Board():

    # list of all 8 directions on the board, as (x,y) offsets
    _directions = [(1,1),(1,0),(1,-1),(0,-1),(-1,-1),(-1,0),(-1,1),(0,1)]

    def __init__(self, n):
        "Set up initial board configuration."

        self.n = n
        # Create the empty board array.
        self.pieces = [None]*self.n
        for i in range(self.n):
            self.pieces[i] = [0]*self.n

        # Set up the initial 4 pieces.
        self.pieces[int(self.n/2)-1][int(self.n/2)] = 1
        self.pieces[int(self.n/2)][int(self.n/2)-1] = 1
        self.pieces[int(self.n/2)-1][int(self.n/2)-1] = -1;
        self.pieces[int(self.n/2)][int(self.n/2)] = -1;

    # add [][] indexer syntax to the Board
    def __getitem__(self, index):
        return self.pieces[index]

    def countDiff(self, color):
        """Counts the # pieces of the given color
        (1 for white, -1 for black, 0 for empty spaces)"""
        count = 0
        for y in range(self.n):
            for x in range(self.n):

```

```

        if self[x][y]==color:
            count += 1
        if self[x][y]==-color:
            count -= 1
    return count

def get_legal_moves(self, color):
    """Returns all the legal moves for the given color.
    (1 for white, -1 for black)
    """
    moves = set() # stores the legal moves.

    # Get all the squares with pieces of the given color.
    for y in range(self.n):
        for x in range(self.n):
            if self[x][y]==color:
                newmoves = self.get_moves_for_square((x,y))
                moves.update(newmoves)
    return list(moves)

def has_legal_moves(self, color):
    for y in range(self.n):
        for x in range(self.n):
            if self[x][y]==color:
                newmoves = self.get_moves_for_square((x,y))
                if len(newmoves)>0:
                    return True
    return False

def get_moves_for_square(self, square):
    """Returns all the legal moves that use the given square as a base.
    That is, if the given square is (3,4) and it contains a black piece,
    and (3,5) and (3,6) contain white pieces, and (3,7) is empty, one
    of the returned moves is (3,7) because everything from there to (3,4)
    is flipped.
    """
    (x,y) = square

    # determine the color of the piece.
    color = self[x][y]

    # skip empty source squares.
    if color==0:
        return None

    # search all possible directions.
    moves = []
    for direction in self.__directions:
        move = self._discover_move(square, direction)
        if move:
            # print(square,move,direction)
            moves.append(move)

    # return the generated move list
    return moves

def execute_move(self, move, color):

```

```

    """Perform the given move on the board; flips pieces as necessary.
    color gives the color pf the piece to play (1=white,-1=black)
    """

    #Much like move generation, start at the new piece's square and
    #follow it on all 8 directions to look for a piece allowing flipping

    # Add the piece to the empty square.
    # print(move)
    flips = [flip for direction in self.__directions
              for flip in self._get_flips(move, direction, color)]
    assert len(list(flips))>0
    for x, y in flips:
        #print(self[x][y],color)
        self[x][y] = color

def _discover_move(self, origin, direction):
    """ Returns the endpoint for a legal move, starting at the given ori
    moving by the given increment."""
    x, y = origin
    color = self[x][y]
    flips = []

    for x, y in Board._increment_move(origin, direction, self.n):
        if self[x][y] == 0:
            if flips:
                # print("Found", x,y)
                return (x, y)
            else:
                return None
        elif self[x][y] == color:
            return None
        elif self[x][y] == -color:
            # print("Flip",x,y)
            flips.append((x, y))

def _get_flips(self, origin, direction, color):
    """ Gets the list of flips for a vertex and direction to use with the
    execute_move function """
    #initialize variables
    flips = [origin]

    for x, y in Board._increment_move(origin, direction, self.n):
        #print(x,y)
        if self[x][y] == 0:
            return []
        if self[x][y] == -color:
            flips.append((x, y))
        elif self[x][y] == color and len(flips) > 0:
            #print(flips)
            return flips

    return []

@staticmethod
def _increment_move(move, direction, n):
    # print(move)

```



```

    """ Generator expression for incrementing moves """
    move = list(map(sum, zip(move, direction)))
    #move = (move[0]+direction[0], move[1]+direction[1])
    while all(map(lambda x: 0 <= x < n, move)):
        #while 0<=move[0] and move[0]<n and 0<=move[1] and move[1]<n:
            yield move
            move=list(map(sum,zip(move,direction)))
            #move = (move[0]+direction[0],move[1]+direction[1])

class OthelloGame():
    def __init__(self, n):
        self.n = n

    def getInitBoard(self):
        # return initial board (numpy board)
        b = Board(self.n)
        return np.array(b.pieces)

    def getBoardSize(self):
        # (a,b) tuple
        return (self.n, self.n)

    def getActionSize(self):
        # return number of actions
        return self.n*self.n + 1

    def getNextState(self, board, player, action):
        # if player takes action on board, return next (board,player)
        # action must be a valid move
        if action == self.n*self.n:
            return (board, -player)
        b = Board(self.n)
        b.pieces = np.copy(board)
        move = (int(action/self.n), action%self.n)
        b.execute_move(move, player)
        return (b.pieces, -player)

    def getValidMoves(self, board, player):
        # return a fixed size binary vector
        valids = [0]*self.getActionSize()
        b = Board(self.n)
        b.pieces = np.copy(board)
        legalMoves = b.get_legal_moves(player)
        if len(legalMoves)==0:
            valids[-1]=1
            return np.array(valids)
        for x, y in legalMoves:
            valids[self.n*x+y]=1
        return np.array(valids)

    def getGameEnded(self, board, player):
        # return 0 if not ended, 1 if player 1 won, -1 if player 1 lost
        # player = 1
        b = Board(self.n)
        b.pieces = np.copy(board)
        if b.has_legal_moves(player):
            return 0

```

```

        if b.has_legal_moves(-player):
            return 0
        if b.countDiff(player) > 0:
            return 1
        return -1

def getCanonicalForm(self, board, player):
    # return state if player==1, else return -state if player==-1
    return player*board

def getSymmetries(self, board, pi):
    # mirror, rotational
    assert(len(pi) == self.n**2+1) # 1 for pass
    pi_board = np.reshape(pi[:-1], (self.n, self.n))
    l = []

    for i in range(1, 5):
        for j in [True, False]:
            newB = np.rot90(board, i)
            newPi = np.rot90(pi_board, i)
            if j:
                newB = np.fliplr(newB)
                newPi = np.fliplr(newPi)
            l += [(newB, list(newPi.ravel()) + [pi[-1]])]
    return l

def stringRepresentation(self, board):
    # 8x8 numpy array (canonical board)
    return board.tostring()

def getScore(self, board, player):
    b = Board(self.n)
    b.pieces = np.copy(board)
    return b.countDiff(player)

def display(board):
    n = board.shape[0]

    for y in range(n):
        print (y,"|",end="")
    print("")
    print(" -----")
    for y in range(n):
        print(y, "|",end="") # print the row #
        for x in range(n):
            piece = board[y][x] # get the piece to print
            if piece == -1: print("b ",end="")
            elif piece == 1: print("W ",end="")
            else:
                if x==n:
                    print("-",end="")
                else:
                    print("- ",end="")
        print("|")

    print(" -----")

```

```

class OthelloNNet():
    def __init__(self, game, args):
        # game params
        self.board_x, self.board_y = game.getBoardSize()
        self.action_size = game.getActionSize()
        self.args = args

        # Renaming functions
        Relu = tf.nn.relu
        Tanh = tf.nn.tanh
        BatchNormalization = tf.layers.batch_normalization
        Dropout = tf.layers.dropout
        Dense = tf.layers.dense

        # Neural Net
        self.graph = tf.Graph()
        with self.graph.as_default():
            self.input_boards = tf.placeholder(tf.float32, shape=[None, self.board_x, self.board_y])
            self.dropout = tf.placeholder(tf.float32)
            self.isTraining = tf.placeholder(tf.bool, name="is_training")

            x_image = tf.reshape(self.input_boards, [-1, self.board_x, self.board_y, 1])
            h_conv1 = Relu(BatchNormalization(self.conv2d(x_image, args['num_channels'], kernel_size=[3,3], padding='same'), axis=-1))
            h_conv2 = Relu(BatchNormalization(self.conv2d(h_conv1, args['num_channels'], kernel_size=[3,3], padding='same'), axis=-1))
            h_conv3 = Relu(BatchNormalization(self.conv2d(h_conv2, args['num_channels'], kernel_size=[3,3], padding='same'), axis=-1))
            h_conv4 = Relu(BatchNormalization(self.conv2d(h_conv3, args['num_channels'], kernel_size=[3,3], padding='same'), axis=-1))
            h_conv4_flat = tf.reshape(h_conv4, [-1, args['num_channels']*(self.board_x*self.board_y)])
            s_fc1 = Dropout(Relu(BatchNormalization(Dense(h_conv4_flat, 1024))), self.dropout)
            s_fc2 = Dropout(Relu(BatchNormalization(Dense(s_fc1, 512), axis=-1)), self.dropout)
            self.pi = Dense(s_fc2, self.action_size)
            self.prob = tf.nn.softmax(self.pi)
            self.v = Tanh(Dense(s_fc2, 1))

            self.calculate_loss()

        def conv2d(self, x, out_channels, padding):
            return tf.layers.conv2d(x, out_channels, kernel_size=[3,3], padding=padding, data_format='channels_last')

        def calculate_loss(self):
            self.target_pis = tf.placeholder(tf.float32, shape=[None, self.action_size])
            self.target_vs = tf.placeholder(tf.float32, shape=[None])
            self.loss_pi = tf.losses.softmax_cross_entropy(self.target_pis, self.prob)
            self.loss_v = tf.losses.mean_squared_error(self.target_vs, self.v)
            self.total_loss = self.loss_pi + self.loss_v
            update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
            with tf.control_dependencies(update_ops):
                self.train_step = tf.train.AdamOptimizer(self.args['lr']).minimize(self.total_loss)

class Coach():
    """
    This class executes the self-play + learning. It uses the functions defined in Game and NeuralNet. args are specified in main.py.
    """
    def __init__(self, game, nnet, args):
        self.game = game
        self.board = game.getInitBoard()
        self.nnet = nnet

```

```

self.pnet = self.nnet.__class__(self.game) # the competitor network
self.args = args
self.mcts = MCTS(self.game, self.nnet, self.args)

def executeEpisode(self):
    """
    This function executes one episode of self-play, starting with player 1.
    As the game is played, each turn is added as a training example to
    trainExamples. The game is played till the game ends. After the game
    ends, the outcome of the game is used to assign values to each example
    in trainExamples.

    It uses a temp=1 if episodeStep < tempThreshold, and thereafter
    uses temp=0.

    Returns:
        trainExamples: a list of examples of the form (canonicalBoard, pi, v)
        pi is the MCTS informed policy vector, v is +1 if
        the player eventually won the game, else -1.
    """
    trainExamples = []
    self.board = self.game.getInitBoard()
    self.curPlayer = 1
    episodeStep = 0

    while True:
        episodeStep += 1
        canonicalBoard = self.game.getCanonicalForm(self.board, self.curPlayer)
        temp = int(episodeStep < self.args['tempThreshold'])

        pi = self.mcts.getActionProb(canonicalBoard, temp=temp)
        sym = self.game.getSymmetries(canonicalBoard, pi)
        for b, p in sym:
            trainExamples.append([b, self.curPlayer, p, None])

        action = np.random.choice(len(pi), p=pi)
        self.board, self.curPlayer = self.game.getNextState(self.board, self.curPlayer, action)

        r = self.game.getGameEnded(self.board, self.curPlayer)

        if r!=0:
            return [(x[0], x[2], r*((-1)**(x[1]!=self.curPlayer))) for x in trainExamples]

def learn(self):
    """
    Performs numIters iterations with numEps episodes of self-play in each
    iteration. After every iteration, it retrains neural network with
    examples in trainExamples (which has a maximum length of maxlenOfQueue).
    It then pits the new neural network against the old one and accepts it
    only if it wins >= updateThreshold fraction of games.
    """

    trainExamples = deque([], maxlen=self.args['maxlenOfQueue'])
    for i in range(self.args['numIters']):
        # bookkeeping
        print('-----ITER ' + str(i+1) + '-----')
        #
        end = time.time()

```

```

        for eps in range(self.args['numEps']):
            self.mcts = MCTS(self.game, self.nnet, self.args)    # reset
            trainExamples += self.executeEpisode()

            # bookkeeping + plot progress
            eps_time.update(time.time() - end)
            end = time.time()

        # training new network, keeping a copy of the old one
        self.nnet.save_checkpoint(folder=self.args['checkpoint'], filename=self.args['checkpoint'] + '.nnet')
        self.pnet.load_checkpoint(folder=self.args['checkpoint'], filename=self.args['checkpoint'] + '.pnet')
        pmcts = MCTS(self.game, self.pnet, self.args)

        self.nnet.train(trainExamples)
        nmcts = MCTS(self.game, self.nnet, self.args)

        print('PITTING AGAINST PREVIOUS VERSION')
        arenal = Arena(lambda x: np.argmax(pmcts.getActionProb(x, temp=0)),
                       lambda x: np.argmax(nmcts.getActionProb(x, temp=0)))
        print('before playgames')
        pwins, nwins, draws = arenal.playGames(self.args['arenaCompare'])

        print('NEW/PREV WINS : %d / %d ; DRAWS : %d' % (nwins, pwins, draws))
        if pwins+nwins > 0 and float(nwins)/(pwins+nwins) < self.args['updateThreshold']:
            print('REJECTING NEW MODEL')
            self.nnet.load_checkpoint(folder=self.args['checkpoint'], filename=self.args['checkpoint'] + '.nnet')
        else:
            print('ACCEPTING NEW MODEL')
            self.nnet.save_checkpoint(folder=self.args['checkpoint'], filename=self.args['checkpoint'] + '.nnet')
            self.pnet.save_checkpoint(folder=self.args['checkpoint'], filename=self.args['checkpoint'] + '.pnet')

class NNetWrapper():
    def __init__(self, game):
        self.nnet = OthelloNNet(game, args)
        self.board_x, self.board_y = game.getBoardSize()
        self.action_size = game.getActionSize()

        self.sess = tf.Session(graph=self.nnet.graph)
        self.saver = None
        with tf.Session() as temp_sess:
            temp_sess.run(tf.global_variables_initializer())
            self.sess.run(tf.variables_initializer(self.nnet.graph.get_collection(tf.GraphKeys.GLOBAL_VARIABLES)))

    def train(self, examples):
        """
        examples: list of examples, each example is of form (board, pi, v)
        """

        for epoch in range(args['epochs']):
            print('EPOCH ::: ' + str(epoch+1))
            # data_time = AverageMeter()
            # batch_time = AverageMeter()
            # pi_losses = AverageMeter()
            # v_losses = AverageMeter()
            end = time.time()

```

```

        batch_idx = 0

        # self.sess.run(tf.local_variables_initializer())
        while batch_idx < int(len(examples)/args['batch_size']):
            sample_ids = np.random.randint(len(examples), size=args['batch_size'])
            boards, pis, vs = list(zip(*[examples[i] for i in sample_ids]))

            # predict and compute gradient and do SGD step
            input_dict = {self.nnet.input_boards: boards, self.nnet.target_pis: pis, self.nnet.target_vs: vs}

            # measure data loading time
            data_time.update(time.time() - end)

            # record loss
            self.sess.run(self.nnet.train_step, feed_dict=input_dict)
            pi_loss, v_loss = self.sess.run([self.nnet.loss_pi, self.nnet.loss_v], feed_dict=input_dict)
            pi_losses.update(pi_loss, len(boards))
            v_losses.update(v_loss, len(boards))

            # measure elapsed time
            batch_time.update(time.time() - end)
            end = time.time()
            batch_idx += 1

            # plot progress

def predict(self, board):
    """
    board: np array with board
    """
    # timing
    start = time.time()

    # preparing input
    board = board[np.newaxis, :, :]

    # run
    prob, v = self.sess.run([self.nnet.prob, self.nnet.v], feed_dict={self.nnet.input_boards: board})

    #print('PREDICTION TIME TAKEN : {0:03f}'.format(time.time()-start))
    return prob[0], v[0]

def save_checkpoint(self, folder='checkpoint', filename='checkpoint.pth.tar'):
    filepath = os.path.join(folder, filename)
    if not os.path.exists(folder):
        print("Checkpoint Directory does not exist! Making directory {}".format(folder))
        os.mkdir(folder)
    else:
        print("Checkpoint Directory exists! ")
    if self.saver == None:
        self.saver = tf.train.Saver(self.nnet.graph.get_collection('variables'))
    with self.nnet.graph.as_default():
        self.saver.save(self.sess, filepath)

def load_checkpoint(self, folder='checkpoint', filename='checkpoint.pth.tar'):
    filepath = os.path.join(folder, filename)

```

```

        if not os.path.exists(filepath+'.meta'):
            raise("No model in path {}".format(filepath))
        with self.nnet.graph.as_default():
            self.saver = tf.train.Saver()
            self.saver.restore(self.sess, filepath)

if __name__=="__main__":
    g = OthelloGame(6)
    nnet = NNetWrapper(g)

    if args['load_model']:
        nnet.load_checkpoint(args['load_folder_file[0]'], args['load_folder_

    c = Coach(g, nnet, args)
    c.learn()

```

```

/Users/monilshah/anaconda3/envs/tensorflow/lib/python3.5/importlib/_bootstrap.py:222: RuntimeWarning: compiletime version 3.6 of module 'tensorflow.python.framework.fast_tensor_util' does not match runtime version 3.5
    return f(*args, **kwargs)

```

```

WARNING:tensorflow:From /Users/monilshah/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/python/ops/losses/losses_impl.py:691: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

```

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

```

-----ITER 1-----

```

Checkpoint Directory exists!

```

INFO:tensorflow:Restoring parameters from ./temp/temp.pth.tar

```

```

EPOCH ::: 1

```

```

-----
--

```

```

KeyboardInterrupt                                Traceback (most recent call last)

```

```

<ipython-input-1-99e7e4c0ad5d> in <module>()

```

```

    751

```

```

    752     c = Coach(g, nnet, args)

```

```

--> 753     c.learn()

```

```

<ipython-input-1-99e7e4c0ad5d> in learn(self)

```

```

    638         pmcts = MCTS(self.game, self.pnet, self.args)

```

```

    639

```

```

--> 640         self.nnet.train(trainExamples)

```

```

    641         nmcts = MCTS(self.game, self.nnet, self.args)

```

```

    642

```

```

<ipython-input-1-99e7e4c0ad5d> in train(self, examples)

```

```

    694

```

```

    695         # record loss

```

```

--> 696         self.sess.run(self.nnet.train_step, feed_dict=input_dict)

```

```

697 #                pi_loss, v_loss = self.sess.run([self.nnet.loss
_pi, self.nnet.loss_v], feed_dict=input_dict)
698 #                pi_losses.update(pi_loss, len(boards))

~/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/pytho
n/client/session.py in run(self, fetches, feed_dict, options, run_metadat
a)
    893     try:
    894         result = self._run(None, fetches, feed_dict, options_ptr,
--> 895                             run_metadata_ptr)
    896     if run_metadata:
    897         proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

~/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/pytho
n/client/session.py in _run(self, handle, fetches, feed_dict, options, ru
n_metadata)
    1126     if final_fetches or final_targets or (handle and feed_dict_ten
nsor):
    1127         results = self._do_run(handle, final_targets, final_fetche
s,
-> 1128                             feed_dict_tensor, options, run_metad
ata)
    1129     else:
    1130         results = []

~/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/pytho
n/client/session.py in _do_run(self, handle, target_list, fetch_list, fee
d_dict, options, run_metadata)
    1342     if handle is None:
    1343         return self._do_call(_run_fn, self._session, feeds, fetches,
s, targets,
-> 1344                             options, run_metadata)
    1345     else:
    1346         return self._do_call(_prun_fn, self._session, handle, feeds
, fetches)

~/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/pytho
n/client/session.py in _do_call(self, fn, *args)
    1348     def _do_call(self, fn, *args):
    1349         try:
-> 1350             return fn(*args)
    1351     except errors.OpError as e:
    1352         message = compat.as_text(e.message)

~/anaconda3/envs/tensorflow/lib/python3.5/site-packages/tensorflow/pytho
n/client/session.py in _run_fn(session, feed_dict, fetch_list, target_lis
t, options, run_metadata)
    1327         return tf_session.TF_Run(session, options,
    1328                                     feed_dict, fetch_list, target_
list,
-> 1329                                     status, run_metadata)
    1330
    1331     def _prun_fn(session, handle, feed_dict, fetch_list):

```

KeyboardInterrupt:

In []: