

```
In [1]: # Make TFGAN models and TF-Slim models discoverable.  
import sys  
import os  
# This is needed since the notebook is stored in the `tensorflow/models/gan`  
sys.path.append('.')  
sys.path.append(os.path.join('.', 'slim'))
```

```
In [2]: from __future__ import absolute_import  
from __future__ import division  
from __future__ import print_function  
  
%matplotlib inline  
import matplotlib.pyplot as plt  
import numpy as np  
import time  
import functools  
from six.moves import xrange # pylint: disable=redefined-builtin  
  
import tensorflow as tf  
  
# Main TFGAN library.  
tfgan = tf.contrib.gan  
  
# TFGAN MNIST examples from `tensorflow/models`.  
from mnist import data_provider  
from mnist import util  
  
# TF-Slim data provider.  
from datasets import download_and_convert_mnist  
  
# Shortcuts for later.  
queues = tf.contrib.slim.queues  
layers = tf.contrib.layers  
ds = tf.contrib.distributions  
framework = tf.contrib.framework
```

```
In [3]: leaky_relu = lambda net: tf.nn.leaky_relu(net, alpha=0.01)
```

```
def visualize_training_generator(train_step_num, start_time, data_np):
    """Visualize generator outputs during training.

    Args:
        train_step_num: The training step number. A python integer.
        start_time: Time when training started. The output of `time.time()`.
            python float.
        data: Data to plot. A numpy array, most likely from an evaluated Tensor
            tensor.
    """
    print('Training step: %i' % train_step_num)
    time_since_start = (time.time() - start_time) / 60.0
    print('Time since start: %f m' % time_since_start)
    print('Steps per min: %f' % (train_step_num / time_since_start))
    plt.axis('off')
    plt.imshow(np.squeeze(data_np), cmap='gray')
    plt.show()

def visualize_digits(tensor_to_visualize):
    """Visualize an image once. Used to visualize generator before training.

    Args:
        tensor_to_visualize: An image tensor to visualize. A python Tensor.
    """
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        with queues.QueueRunners(sess):
            images_np = sess.run(tensor_to_visualize)
    plt.axis('off')
    plt.imshow(np.squeeze(images_np), cmap='gray')

def evaluate_tfgan_loss(gan_loss, name=None):
    """Evaluate GAN losses. Used to check that the graph is correct.

    Args:
        gan_loss: A GANLoss tuple.
        name: Optional. If present, append to debug output.
    """
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        with queues.QueueRunners(sess):
            gen_loss_np = sess.run(gan_loss.generator_loss)
            dis_loss_np = sess.run(gan_loss.discriminator_loss)
    if name:
        print('%s generator loss: %f' % (name, gen_loss_np))
        print('%s discriminator loss: %f' % (name, dis_loss_np))
    else:
        print('Generator loss: %f' % gen_loss_np)
        print('Discriminator loss: %f' % dis_loss_np)
```


```
In [4]: MNIST_DATA_DIR = '/tmp/mnist-data'
```

```
if not tf.gfile.Exists(MNIST_DATA_DIR):  
    tf.gfile.MakeDirs(MNIST_DATA_DIR)  
  
download_and_convert_mnist.run(MNIST_DATA_DIR)
```

Dataset files already exist. Exiting without re-creating them.

```
In [5]: tf.reset_default_graph()
```

```
# Define our input pipeline. Pin it to the CPU so that the GPU can be reserved  
# for forward and backwards propagation.  
batch_size = 32  
with tf.device('/cpu:0'):  
    real_images, _, _ = data_provider.provide_data(  
        'train', batch_size, MNIST_DATA_DIR)  
  
# Sanity check that we're getting images.  
check_real_digits = tfgan.eval.image_resaper(  
    real_images[:20,...], num_cols=10)  
visualize_digits(check_real_digits)
```



```
In [6]: def generator_fn(noise, weight_decay=2.5e-5, is_training=True):
        """Simple generator to produce MNIST images.

        Args:
            noise: A single Tensor representing noise.
            weight_decay: The value of the l2 weight decay.
            is_training: If `True`, batch norm uses batch statistics. If `False`
                        norm uses the exponential moving average collected from population
                        statistics.

        Returns:
            A generated image in the range [-1, 1].
        """
        with framework.arg_scope(
            [layers.fully_connected, layers.conv2d_transpose],
            activation_fn=tf.nn.relu, normalizer_fn=layers.batch_norm,
            weights_regularizer=layers.l2_regularizer(weight_decay)),\
            framework.arg_scope([layers.batch_norm], is_training=is_training,
                                zero_debias_moving_mean=True):
            net = layers.fully_connected(noise, 1024)
            net = layers.fully_connected(net, 7 * 7 * 256)
            net = tf.reshape(net, [-1, 7, 7, 256])
            net = layers.conv2d_transpose(net, 64, [4, 4], stride=2)
            net = layers.conv2d_transpose(net, 32, [4, 4], stride=2)
            # Make sure that generator output is in the same range as `inputs`
            # ie [-1, 1].
            net = layers.conv2d(net, 1, 4, normalizer_fn=None, activation_fn=tf.nn.tanh)

        return net
```

```
In [7]: def discriminator_fn(img, unused_conditioning, weight_decay=2.5e-5,
        is_training=True):
        """Discriminator network on MNIST digits.

        Args:
            img: Real or generated MNIST digits. Should be in the range [-1, 1].
            unused_conditioning: The TFGAN API can help with conditional GANs, which
                would require extra `condition` information to both the generator and
                discriminator. Since this example is not conditional, we do not use this
                argument.
            weight_decay: The L2 weight decay.
            is_training: If `True`, batch norm uses batch statistics. If `False`, batch
                norm uses the exponential moving average collected from population batch
                statistics.

        Returns:
            Logits for the probability that the image is real.
        """
        with framework.arg_scope([layers.conv2d, layers.fully_connected],
                                activation_fn=leaky_relu, normalizer_fn=None,
                                weights_regularizer=layers.l2_regularizer(weight_decay),
                                biases_regularizer=layers.l2_regularizer(weight_decay)):
            net = layers.conv2d(img, 64, [4, 4], stride=2)
            net = layers.conv2d(net, 128, [4, 4], stride=2)
            net = layers.flatten(net)
            with framework.arg_scope([layers.batch_norm], is_training=is_training):
                net = layers.fully_connected(net, 1024, normalizer_fn=layers.batch_norm)
            return layers.linear(net, 1)
```

```
In [8]: noise_dims = 64
gan_model = tfgan.gan_model(
    generator_fn,
    discriminator_fn,
    real_data=real_images,
    generator_inputs=tf.random_normal([batch_size, noise_dims]))

# Sanity check that generated images before training are garbage.
check_generated_digits = tfgan.eval.image_resampler(
    gan_model.generated_data[:20, ...], num_cols=10)
visualize_digits(check_generated_digits)
```



```

In [9]: # We can use the minimax loss from the original paper.
vanilla_gan_loss = tfgan.gan_loss(
    gan_model,
    generator_loss_fn=tfgan.losses.minimax_generator_loss,
    discriminator_loss_fn=tfgan.losses.minimax_discriminator_loss)

# We can use the Wasserstein loss (https://arxiv.org/abs/1701.07875) with the
# gradient penalty from the improved Wasserstein loss paper
# (https://arxiv.org/abs/1704.00028).
improved_wgan_loss = tfgan.gan_loss(
    gan_model,
    # We make the loss explicit for demonstration, even though the default is
    # Wasserstein loss.
    generator_loss_fn=tfgan.losses.wasserstein_generator_loss,
    discriminator_loss_fn=tfgan.losses.wasserstein_discriminator_loss,
    gradient_penalty_weight=1.0)

# We can also define custom losses to use with the rest of the TFGAN framework.
def silly_custom_generator_loss(gan_model, add_summaries=False):
    return tf.reduce_mean(gan_model.discriminator_gen_outputs)
def silly_custom_discriminator_loss(gan_model, add_summaries=False):
    return (tf.reduce_mean(gan_model.discriminator_gen_outputs) -
            tf.reduce_mean(gan_model.discriminator_real_outputs))
custom_gan_loss = tfgan.gan_loss(
    gan_model,
    generator_loss_fn=silly_custom_generator_loss,
    discriminator_loss_fn=silly_custom_discriminator_loss)

# Sanity check that we can evaluate our losses.
for gan_loss, name in [(vanilla_gan_loss, 'vanilla loss'),
                       (improved_wgan_loss, 'improved wgan loss'),
                       (custom_gan_loss, 'custom loss')]:
    evaluate_tfgan_loss(gan_loss, name)

vanilla loss generator loss: -1.450439
vanilla loss discriminator loss: 1.674405
improved wgan loss generator loss: 0.539317
improved wgan loss discriminator loss: -0.020071
custom loss generator loss: -0.347576
custom loss discriminator loss: 0.024440

```

```

In [10]: generator_optimizer = tf.train.AdamOptimizer(0.001, beta1=0.5)
discriminator_optimizer = tf.train.AdamOptimizer(0.0001, beta1=0.5)
gan_train_ops = tfgan.gan_train_ops(
    gan_model,
    improved_wgan_loss,
    generator_optimizer,
    discriminator_optimizer)

```

```
In [11]: num_images_to_eval = 500
MNIST_CLASSIFIER_FROZEN_GRAPH = './mnist/data/classify_mnist_graph_def.pb'

# For variables to load, use the same variable scope as in the train job.
with tf.variable_scope('Generator', reuse=True):
    eval_images = gan_model.generator_fn(
        tf.random_normal([num_images_to_eval, noise_dims]),
        is_training=False)

# Calculate Inception score.
eval_score = util.mnist_score(eval_images, MNIST_CLASSIFIER_FROZEN_GRAPH)

# Calculate Frechet Inception distance.
with tf.device('/cpu:0'):
    real_images, _, _ = data_provider.provide_data(
        'train', num_images_to_eval, MNIST_DATA_DIR)
    frechet_distance = util.mnist_frechet_distance(
        real_images, eval_images, MNIST_CLASSIFIER_FROZEN_GRAPH)


# Reshape eval images for viewing.
generated_data_to_visualize = tfgan.eval.image_resaper(
    eval_images[:20,...], num_cols=10)
```

```
In [12]: # We have the option to train the discriminator more than one step for every
# step of the generator. In order to do this, we use a `GANTrainSteps` with
# desired values. For this example, we use the default 1 generator train step
# for every discriminator train step.
train_step_fn = tfgan.get_sequential_train_steps()


global_step = tf.train.get_or_create_global_step()
loss_values, mnist_scores, frechet_distances = [], [], []

with tf.train.SingularMonitoredSession() as sess:
    start_time = time.time()
    for i in xrange(1601):
        cur_loss, _ = train_step_fn(
            sess, gan_train_ops, global_step, train_step_kwargs={})
        loss_values.append((i, cur_loss))
        if i % 200 == 0:
            mnist_score, f_distance, digits_np = sess.run(
                [eval_score, frechet_distance, generated_data_to_visualize])
            mnist_scores.append((i, mnist_score))
            frechet_distances.append((i, f_distance))
            print('Current loss: %f' % cur_loss)
            print('Current MNIST score: %f' % mnist_scores[-1][1])
            print('Current Frechet distance: %f' % frechet_distances[-1][1])
            visualize_training_generator(i, start_time, digits_np)
```

```
Current loss: -0.744440
Current MNIST score: 1.000062
Current Frechet distance: 339.406494
Training step: 0
Time since start: 0.065372 m
Steps per min: 0.000000
```



```
Current loss: -3.240865
Current MNIST score: 2.691993
Current Frechet distance: 166.751053
Training step: 200
Time since start: 1.576845 m
Steps per min: 126.835563
```




```
Current loss: -3.095253
Current MNIST score: 4.904757
Current Frechet distance: 93.779289
Training step: 400
```



Time since start: 3.092882 m
Steps per min: 129.329208




Current loss: -3.088769
Current MNIST score: 6.725487
Current Frechet distance: 39.744904
Training step: 600
Time since start: 4.608689 m
Steps per min: 130.188869



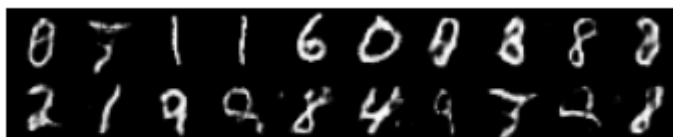
Current loss: -3.074484
Current MNIST score: 7.534400
Current Frechet distance: 19.210173
Training step: 800
Time since start: 6.130447 m
Steps per min: 130.496196



Current loss: -2.957786
Current MNIST score: 7.315704
Current Frechet distance: 24.819916
Training step: 1000
Time since start: 7.653988 m
Steps per min: 130.650852

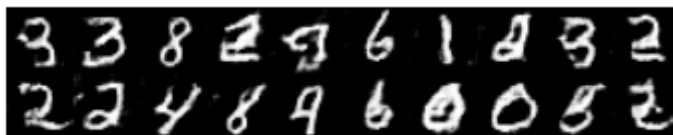


Current loss: -2.981731
Current MNIST score: 6.797893
Current Frechet distance: 53.053493
Training step: 1200
Time since start: 9.164964 m
Steps per min: 130.933413



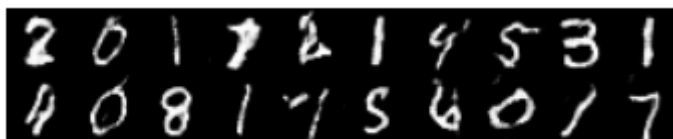
0 5 1 1 6 0 0 8 8 0
2 1 9 9 8 4 9 5 2 8

Current loss: -3.221900
Current MNIST score: 7.398132
Current Frechet distance: 33.662426
Training step: 1400
Time since start: 10.689862 m
Steps per min: 130.965214



3 3 8 2 9 6 1 2 3 2
2 2 4 8 9 6 0 0 5 2

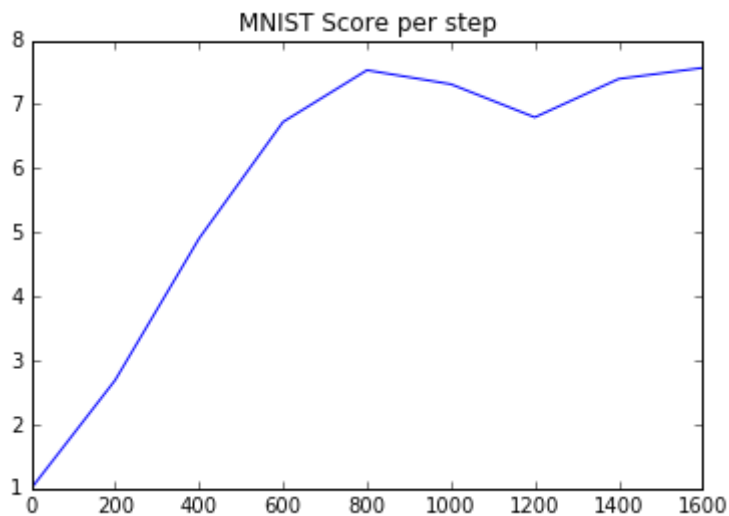
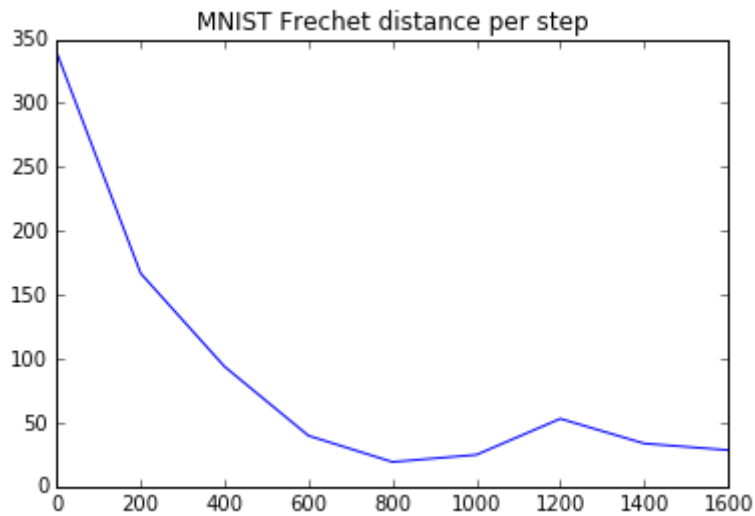
Current loss: -3.512912
Current MNIST score: 7.570933
Current Frechet distance: 28.488590
Training step: 1600
Time since start: 12.209630 m
Steps per min: 131.044103




2 0 1 7 2 1 9 5 3 1
4 0 8 1 7 5 6 0 1 7

```
In [13]: # Plot the eval metrics over time.
plt.title('MNIST Frechet distance per step')
plt.plot(*zip(*frechet_distances))
plt.figure()
plt.title('MNIST Score per step')
plt.plot(*zip(*mnist_scores))
plt.figure()
plt.title('Training loss per step')
plt.plot(*zip(*loss_values))
```

Out[13]: [





In []: